

# Optimizing Pipeline Parallelism for Deep Learning with Activation Checkpointing

Min-Yen Chiang<sup>1</sup>, Tzu-Hsien Tsai<sup>1</sup>, Ding-Yong Hong<sup>2</sup>,  
Pangfeng Liu<sup>1</sup>, Jan-Jan Wu<sup>2</sup>

<sup>1</sup>Department of Computer Science and Information Engineering,  
National Taiwan University, Taipei, Taiwan

<sup>2</sup>Institute of Information Science, Academia Sinica, Taipei, Taiwan

November 26, 2025

# Introduction

- We increase the throughput of DNN training by better pipeline parallelism and activation checkpointing.

# The Importance of Deep Neural Networks

- Deep neural networks solve problems across various fields.
  - Object detection
  - Text summarization
  - Protein structure prediction
  - Traffic forecast

# DNN Training Procedure

- Let  $F_i$  be the forward pass of layer  $i$ .
- Let  $B_i$  be the backward pass of layer  $i$ .
- Let  $a_i$  be the activation of layer  $i$ .
- Let  $g_i$  be the gradient of layer  $i$ .

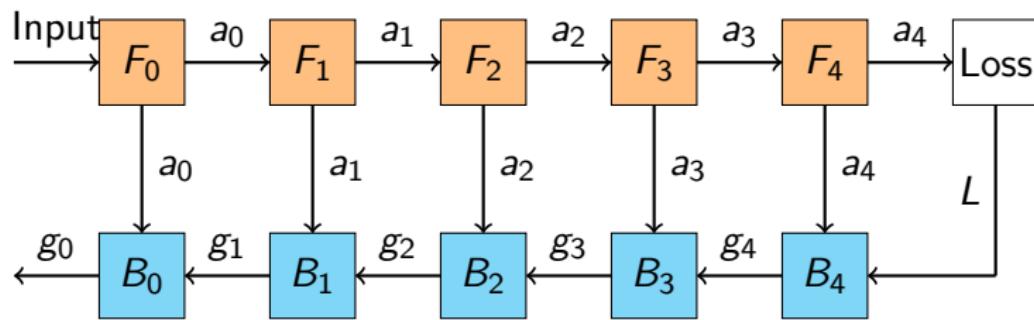


Figure 1: The training procedure of DNNs

# Pipeline Parallelism

- The training procedure of a model is partitioned into multiple stages.
  - Each stage consists of consecutive layers in the model.
  - Each stage is trained on different devices.
- A batch of data is divided into multiple micro batches.

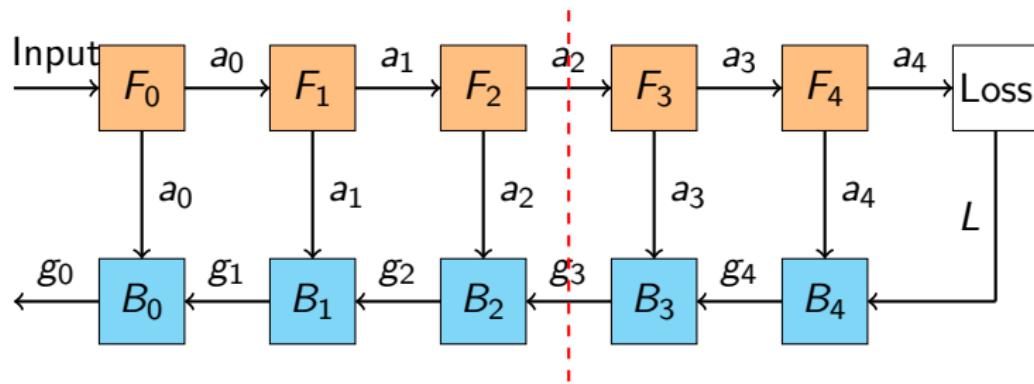


Figure 2: The DNN training procedure using pipeline parallelism

# The Memory Usage of Training a Stage

- The activations occupy most of the memory during DNN training<sup>1</sup>.

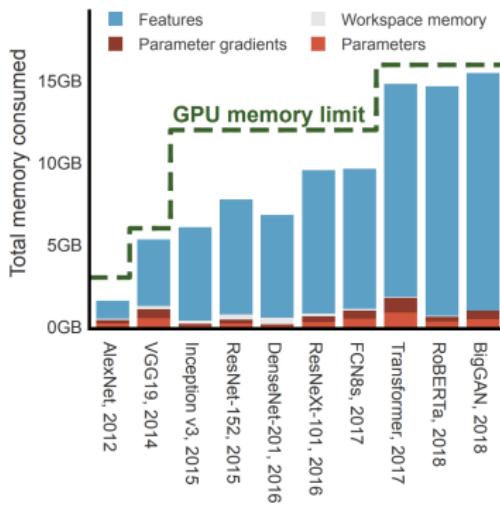


Figure 3: The memory usage of DNN training

<sup>1</sup> Paras Jain et al. "Checkmate: Breaking the memory wall with optimal tensor rematerialization". In: *MLSys.* Vol. 2. 2020, pp. 497–511.

# Activation Checkpointing

- Activation Checkpointing stores a subset of activations during the forward pass and recomputes the discarded activations during the backward pass.
  - The stored activations are called *checkpoints*.

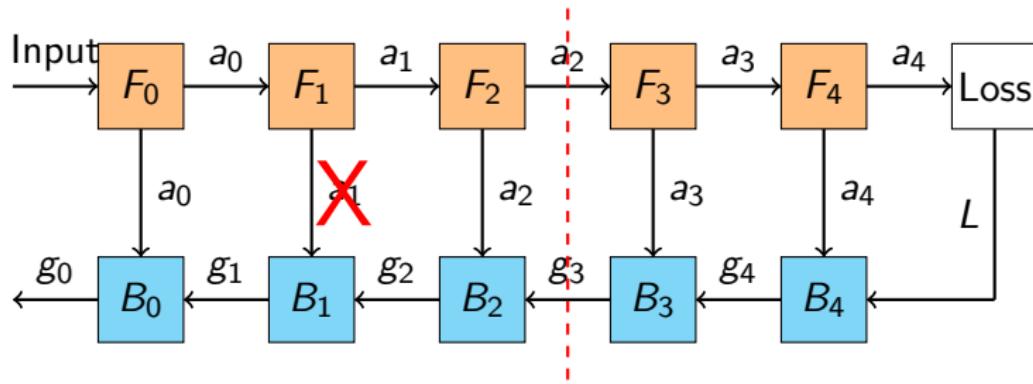


Figure 4: Activation Checkpointing

# Fast-forward Models

- A fast-forward model is the DNN that contains a layer whose output is passed to a non-adjacent layer.
  - The fast-forward edges do not cross each other.
  - Many widely-used DNNs are fast-forward models, including ResNet, GPT-3, and Llama.
- We consider the training procedure of fast-forward models.

# Problems

- How to schedule the micro batches into the pipeline?
- How to partition a model into stages?
- How to select checkpoints in a stage?

# Pipeline Schedule

- We use the one-forward-one-backward (1F1B)<sup>2</sup> pipeline schedule.
- Let  $p$  be the number of stages.
- The 1F1B pipeline schedule passes a micro batch backward whenever the input to this backward pass is ready.
- If there is no such a micro batch on the  $i$ -th stage and the number of micro batches whose activations are stored on the device does not exceed  $p - i + 1$ , then a micro batch is passed forward.

---

<sup>2</sup>Deepak Narayanan et al. “PipeDream: Generalized pipeline parallelism for DNN training”. In: *SOSP*. 2019, pp. 1–15.

# Previous Work on Training Fast-forward Models

- AdaPipe<sup>3</sup> uses a heuristic to partition a fast-forward model.
- AdaPipe uses a heuristic to find the checkpoints in a fast-forward model.

---

<sup>3</sup>Zhenbo Sun et al. “AdaPipe: Optimizing pipeline parallelism with adaptive recomputation and partitioning”. In: *ASPLOS*. 2024, pp. 86–100.

# Contribution

- Our model partition and activation checkpointing algorithms are based on mathematical formulation and dynamic programming.
- Our model partition minimize the maximum stage computation time in the pipeline.
- Our activation checkpointing minimize the stage computation time with memory usage no more than the device memory capacity.
- Our algorithms achieve a higher throughput and a more accurate memory estimation than AdaPipe.

# The Goal of Model Partitioning

- Let  $n$  be the number of layers in a model.
- Let  $h_i$  be the first layer in the  $i$ -th stage.
  - We set  $h_1 = 1$  and  $h_{p+1} = n + 1$ .
- Let the training time of the  $i$ -th stage be  $C_i(h_i, h_{i+1} - 1)$ .
- Our goal is to find the sequence  $(h_1, h_2, \dots, h_p, h_{p+1})$  that minimizes  $S$  in Equation 1.

$$S = \max_{1 \leq i \leq p} C_i(h_i, h_{i+1} - 1) \quad (1)$$

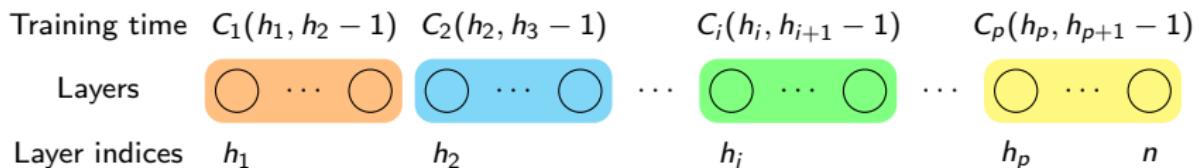


Figure 5: The model partition and stage training time

# The Algorithm for Model Partitioning

- Let  $D(i, q)$  be the shortest maximum stage training time when layers between layer  $i$  and layer  $n$  are partitioned into  $q$  stages.
- The minimum of  $S$  is  $D(1, p)$ .

$$S = \max_{1 \leq i \leq p} C_i(h_i, h_{i+1} - 1)$$

- We derive  $D(i, q)$  by enumerating the last layer of the stage that starts from  $i$ .

$$D(i, q) = \min_{i \leq k < n} \max(C_{p-q+1}(i, k), D(k + 1, q - 1)) \quad (2)$$

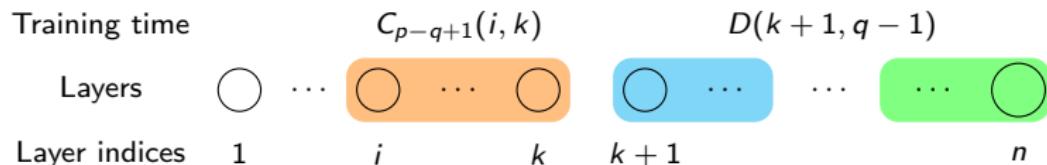


Figure 6: The idea behind Recurrence 2

# PyTorch Constraint

- PyTorch does not support the model partition that splits a fast-forward edge.
- Let  $P$  be the set of layers that are not crossed by a fast-forward edge.

$$D(i, q) = \min_{\substack{i \leq k < n \\ k \in P}} \max(C_{p-q+1}(i, k), D(k + 1, q - 1)) \quad (3)$$

# Stage Training Time and Activation Checkpointing

- Now we construct the stage training time  $C_i(\cdot, \cdot)$ .
- The value of  $C_i(\cdot, \cdot)$  is related to the checkpoint selection.
- We aim to select checkpoints that minimize  $C_i(\cdot, \cdot)$ .
- Each activation is recomputed at most once.
  - This constraint is adopted by many existing checkpointing methods<sup>456</sup>.
- We consider deriving the value of  $C_i(1, n)$ .

---

<sup>4</sup>Tianqi Chen et al. *Training deep nets with sublinear memory cost*. arXiv preprint arXiv:1604.06174. 2016.

<sup>5</sup>Jianwei Feng and Dong Huang. “Optimal gradient checkpoint search for arbitrary computation graphs”. In: CVPR. 2021, pp. 11433–11442.

<sup>6</sup>Ding-Yong Hong et al. “GPU memory usage optimization for backward propagation in deep network training”. In: JPDC 199 (2025), p. 105053

# The Goal of Activation Checkpointing

- Let  $f_j$  be the time for passing forward a micro batch through layer  $j$ .
- Let  $b_j$  be the time for passing backward a micro batch through layer  $j$ .
- Let  $A$  be the set of checkpoints.
- The stage training time is the sum of the forward passing time, the backward passing time, and the time for recomputing the discarded activation, as shown in Equation 4.

$$C_i(1, n) = \sum_{1 \leq j \leq n} f_j + \sum_{1 \leq j \leq n} b_j + \left( \sum_{1 \leq j \leq n} f_j - \sum_{j \in A} f_j \right) \quad (4)$$

- Minimizing  $C_i(1, n)$  is equal to maximizing  $C'_i(1, n)$  in Equation 5.

$$C'_i(1, n) = \sum_{j \in A} f_{c_j} \quad (5)$$

# The Constraint of Activation Checkpointing

- The term  $C'_i(1, n)$  is maximized when  $A$  contains every layer in the model.

$$C'_i(1, n) = \sum_{j \in A} f_{c_j}$$

- The memory usage for storing every activations may exceed the device memory capacity.
- Therefore, our goal of activation checkpointing becomes maximizing the checkpoint computation time  $C'_i(1, n)$  with the memory usage no more than the device memory capacity.

# The Memory Usage for Backward Passing

- The memory usage for training a stage consists of two parts.
  - The checkpoints
  - The gradients and the recomputed activations
- Let  $l$  be the maximum number of micro batches whose activations are stored on the device.
- Let  $d_i$  be the size of the activation of layer  $i$ .
- The memory usage of selecting layer  $i$  as a checkpoint is  $l \times d_i$ .
- Let  $s(i, j)$  be the memory usage for storing the gradients and the recomputed activations when passing a micro batch backward from layer  $j$  to layer  $i$ .

# Maximum Checkpoint Computation Time with Memory Constraint

- Let  $T(i, m)$  be the maximum checkpoint computation time of layers between layer  $i$  and layer  $n$  with the memory usage no more than  $m$  and layer  $i$  as a checkpoint.
- The maximum checkpoint computation time is  $T(1, M)$ , where  $M$  is the device memory capacity.
- We derive  $T(i, m)$  by enumerating the first checkpoint after layer  $i$ .

$$T(i, m) = f_i + \max_{\substack{i < j \leq n \\ s(i,j) \leq m - l \times d_i}} T(j, m - l \times d_i) \quad (6)$$

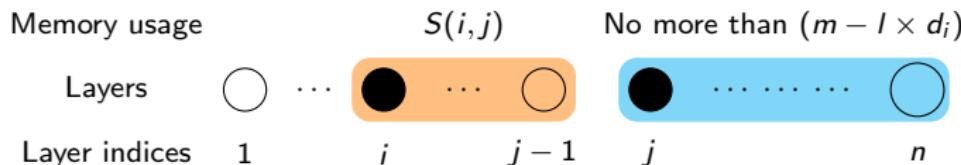


Figure 7: The idea behind Recurrence 7

# PyTorch Constraint for Activation Checkpointing

- PyTorch selects the starting point of a fast-forward edge  $e_f$  as a checkpoint under the following two conditions.
  - Neither the starting point nor the endpoint of  $e_f$  is a checkpoint.
  - The edge  $e_f$  crosses a checkpoint.
- PyTorch applies this mechanism to avoid storing too many activations in the memory
- To ensure that our memory estimation is consistent with the memory usage of PyTorch, our algorithm adopts the PyTorch constraint.

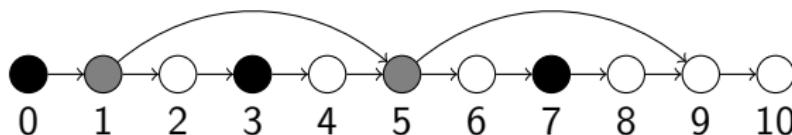


Figure 8: Checkpoints added by PyTorch

# Maximum Checkpoint Computation Time with Memory and PyTorch Constraints

- Let  $u(i)$  be the first layer after layer  $i$  that is the starting point or the endpoint of a fast-forward edge.
- When  $j > u(i)$  and layer  $j$  is crossed by a fast-forward edge  $e_f$ , our algorithm selects  $u(j)$  as the checkpoint.

$$T(i, m) = f_i + \max_{\substack{i < j \leq n \\ s(i, j) \leq m - l \times d_i}} T(j, m - l \times d_i) \quad (7)$$

- The fast-forward edges do not cross each other.

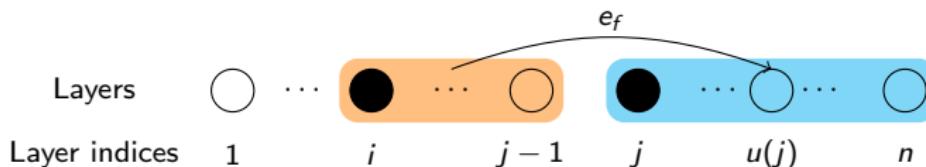


Figure 9: The situation that may violate the PyTorch constraint

# Maximum Checkpoint Computation Time with Memory and PyTorch Constraints

- The function  $T^*(i, m)$  is  $T(i, m)$  with an additional constraint that  $u(i)$  is the checkpoint.
- Let  $U$  be the set consisting of all layers that may violate the PyTorch constraint.

$$T(i, m) = f_i + \max_{\substack{i < j \leq n \\ s(i, j) \leq m - l \times d_i}} \begin{cases} T(j, m - l \times d_i) & j \notin U \\ T^*(j, m - l \times d_i) & j \in U \end{cases} \quad (8)$$

$$T^*(i, m) = f_i + \max_{\substack{i < j \leq u(i) \\ s(i, j) \leq m - l \times d_i}} \begin{cases} T^*(j, m - l \times d_i) & j < u(i) \\ T(j, m - l \times d_i) & j = u(i) \end{cases} \quad (9)$$

# Environment

- We conduct our experiments on the server with 32 2.1 GHz Intel(R) Xeon(R) Gold 6530 processors and four Nvidia A6000 GPUs.
- We use PyTorch to implement our experiments.

# Memory Prediction

- The average difference of our memory estimation and the memory usage reported by PyTorch is 0.9%.

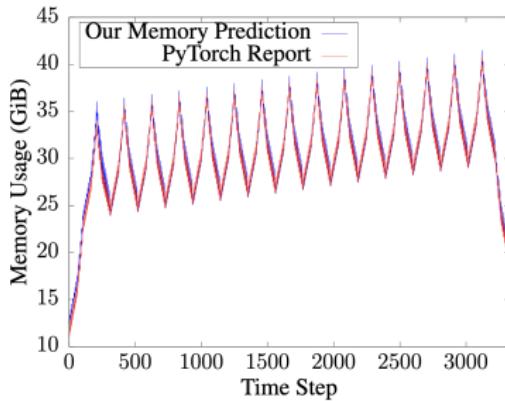


Figure 10: Our memory prediction and the memory usage of PyTorch implementation

# Throughput

- Our algorithm outperforms the state-of-the-art algorithm AdaPipe<sup>7</sup> on models with various sizes and structures.

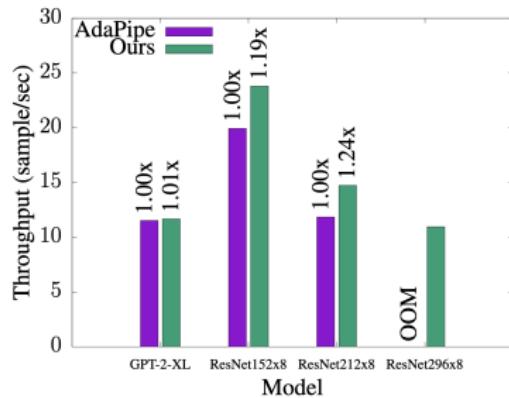


Figure 11: The training throughput of our algorithm and that of AdaPipe

<sup>7</sup>Sun et al., "AdaPipe: Optimizing pipeline parallelism with adaptive recomputation and partitioning".

# Conclusion

- We increase the training throughput of fast-forward models by better pipeline parallelism and activation checkpointing.
- Our model partition minimizes the maximum stage training time.
- Our checkpoint selection minimizes the stage training time under the memory capacity.
- Our methods follow the constraint of PyTorch and thus estimate the memory usage accurately.

Q & A

Q & A